
What Moves the Ceiling on Automated Code Audit: the Auditor, its Rulebook and the Specification, Measured One at a Time

Zhaohe Dong¹

Abstract

An automated auditor reads code it did not write and decides whether to block it. We measure what moves the limit on how much it catches, against ground truth no model supplies (hidden test suites run by an interpreter), varying one factor at a time on 110 defective and 150 correct increments. Eight independent readings of a shipped cross-vendor auditor reach 30.0% union recall [20.0, 40.7] at 16.0% false positives, and the eighth still adds 1.9 points: returns diminish, with no plateau shown. A stronger same-vendor model moves recall by -26.4 points [-37.3, -15.6], a preregistered primary, while flagging 3.3% of correct code against 16.0%; the sign reverses under a weaker severity rule, so the contrast depends on the counting rule and is not a measure of auditing ability. One added rulebook sentence moves flags by +26.8 points [13.6, 40.4] on an exploratory arm, at +12.5 points [-3.6, +26.8] on correct code, and does not measurably move hidden-suite pass rate. Of 57 defects no reading flagged, two raters, one an author, judge 44 to be failures the specification never determined; a model not among the families that missed them, asked a narrower question, labels 38 undetermined. Adding a behavioural rule to the specification with the code held fixed raised adjudicated diagnosis from 0 to 9 of 32 selected instances (+28.1 points [7.7, 51.4]), against 0 under a placebo edit, only where the added rule agreed with the hidden suite; the exact cluster test gives $p = 0.06$. A recall figure is therefore a property of auditor, rulebook, decision rule and benchmark together, and is not portable without its false-positive rate. Every study that made model calls was preregistered and the two unregistered analyses are named; each report is marked by whether it passed independent cross-vendor review, and the reviews are archived beside the

records.

1. Introduction

An automated auditor reads code it did not write and says whether to block it. We ask how much such an auditor catches, and what decides the limit.

That question has parts, though our design measures each one and does not apportion between them. A reviewer's score on a benchmark mixes at least four things together: whether the model can see a defect, whether it judges the defect worth blocking on, whether the benchmark's hidden tests define the behaviour it would have to see, and how many times it is allowed to look. Work on code-review and critic models improves the first two and reports the aggregate gain. We hold the auditor, its instructions and its ground truth fixed, and vary the other three one at a time, against a truth no model supplies: a hidden test suite executed by an interpreter.

Returns to reading again diminish. Eight independent readings of the same increment by the shipped cross-vendor auditor reach 30.0% union recall on defective increments, with a 95% problem-cluster interval of 20.0 to 40.7, at 16.0% [10.1, 22.3] on correct ones. The eighth reading still adds 1.9 points [1.1, 2.8], above the 1.0-point flattening bar we registered in advance, so the curve had not flattened at the budget we could afford and we do not claim saturation. Twenty readings pooled across three auditor families reach 48.2%, at 36.0% false positives on correct code.

Which model reads, and what it is told to look for, move the limit further. Swapping in a stronger model of the generator's own vendor changes union recall by -26.4 points [-37.3, -15.6], a preregistered primary. That contrast does not survive its own controls: the stronger model flags correct code at 3.3% against 16.0%, the sign reverses under a weaker severity rule, and the two routes are sampled differently, one at temperature 0 and one with no sampling parameter sent, so their unions are not equally free to grow. A third model reaches higher observed recall than the shipped auditor at fewer false positives with half the readings, though no paired contrast establishes both at once. Adding one

¹Independent researcher. Correspondence to: Zhaohe Dong <dongzhaohe321418@gmail.com>.

sentence to the auditor’s rulebook, telling it to look for what the visible tests do not cover, moves flags by +26.8 points [13.6, 40.4] on a 56-instance exploratory arm, at 12.5 points more flags on correct code — a cost whose admissible interval is $[-3.6, +26.8]$ and does not exclude zero; under the rule that stratum is flagged at 10 of 56, above the product’s 6.7% false-positive constraint as an observed rate, against 3 of 56 without it; on the contrast that isolates the sentence, hidden-suite pass rate moves +5.4 points $[-0.9, 12.1]$ and does not exclude zero.

Much of what the auditor misses turns out to be failure the specification never determined. We re-rated the never-flagged residual under a rule that asks whether the prose settles the hidden suite’s expected value before it asks whether the visible tests exercised the input. Two raters agree that 44 of 57 residual instances are cases where the prose does not settle it, reversing the same instances’ earlier classification. Where an outside party had already flagged one of these specifications as defective, for their own purposes and before we looked, our raters agree on 9 of the 11 tasks that overlap.

We tried to test that prospectively, by constructing a population whose specification-determinedness would hold by construction, and the attempt failed in a way worth reporting. Independent review established that our construction never enforced the property: the filter registered to recover the failing input is a truthiness check on a field the injector wrote — of 281 recorded witnesses, 179 are strings and 102 are lists, and none is executed — and four of the other five filters accept cases their registered wording excludes. What the population actually contains is small injected edits that survive a sparse visible suite and fail a hidden one. A preregistered probe separates those edits from natural code at 96.7% on the items it answered, *by a feature it does not identify*: its two arms share no problem, so task composition alone could produce the result. The prospective test therefore does not exist yet, and the classification above remains post hoc.

A union recall figure is contingent on the false-positive rate it cost, so the ceiling we measure describes an operating point rather than a constant. We argue that as a methodological point and not from a second measurement: we attempted one on a substrate with much longer specifications, and it closed without a quotable number after three runs (§3.5). Its clearest lesson is the one we do not get to quote — across three extractions of the same protocol, the quantity we had most wanted to use, whether the two substrates’ false-positive intervals meet, came out apart, overlapping, and apart again.

1.1. Preregistration and review

Every study that made model calls was preregistered before its first one, with hypotheses, intervals, kill conditions and stopping rules fixed in advance, and every departure recorded as a numbered amendment that says what was known at the time. Two analyses were not registered in advance, and the paper names both: a re-grading of archived findings, which was never registered, and a third rating whose registration came after its result and is withdrawn. Four preregistered hypotheses were killed by their own rules, two of them our own conjectures, and one of those kills requires a headline in this paper to be stated as a flag rate rather than a recall. Reports are reviewed by a model from a different vendor with access to the raw records, several needing many rounds, and the reviews are archived verbatim beside the results they judge. Every study reported here has completed that review, at between three and twenty-one rounds each, and Section 3.5 collects the seven that finished after this paper was first drafted. We record this as method, and it is what lets a reader read the negative results here as measurements.

2. Method

2.1. Substrate, population and ground truth

Every result in this paper is measured on one substrate: HumanEval+ and MBPP+ (EvalPlus), Python functions whose median task is 41 words of prose and 6 lines of reference code, each with an executable hidden test suite, so every stratum, recall denominator and false-positive population is decided by an interpreter and no model judges a recall or false-positive outcome. The one outcome defined by adjudication, correct diagnosis in the clarification study, is judged jointly by an author and a model, and is labelled so where it appears. The benchmark’s base suite is shown to the generator as the *visible* suite and its plus suite is held back as the *hidden* one. A second substrate, the standard-library half of BigCodeBench, was built and audited and closed without a quotable number, when its preregistered criterion was not met on a third run; nothing from it is reported here.

How the population was built. The corpus has 542 problems; two, HumanEval/32 and Mbpp/590, were excluded by name before generation because their own reference solutions fail their own hidden suites in this harness, which leaves 540. One frozen generator, `claude-haiku-4-5-20251001`, was shown each specification and its visible suite and wrote one solution per problem. A candidate that passes the visible suite and fails the hidden one is a **defective** instance; one that passes both we call **correct**; one that fails the visible suite is discarded from every recall and false-positive population. The defect stratum is small — about 10% of solutions — so, under a

rule registered before the second pass and depending only on execution outcomes, generation was repeated over the same 540 problems until there were at least 100 defective instances or three passes. It stopped at two: 56 and 54 defective, 455 and 455 correct, 29 and 31 discarded. An instance is therefore a (pass, problem) pair, and one problem can contribute two. Recall is measured on all **110** defective instances, which come from **56 problems**; false positives on **150** of the 910 correct instances, drawn once from a fixed seed before any auditor ran, from 143 problems. Seven of the 110 defective instances fail by not terminating rather than by a failed assertion, and the registered population keeps them.

Two things follow from those definitions and are meant literally throughout. *Correct* means *passes the hidden suite*: it is a benchmark label produced by an interpreter, not a guarantee that the program meets its specification, and a suite can be incomplete or wrong. Since this paper’s own third result is that a specification often fails to determine what the hidden suite asks, we cannot also treat passing that suite as correctness in the ordinary sense, and a *false positive* is likewise a flag on test-passing code rather than a demonstrated mistake by the auditor. And the defect population is not a sample of model-written code: it is the **residue that already survived the visible tests**, which is by construction the harder part. A recall figure here is the fraction of that residue an auditor catches, never the fraction of all defects a model writes.

2.2. The auditor and the estimand

The auditor reads the specification, the visible tests and the candidate, and returns findings with severities; an increment is blocked when any finding is graded BLOCKER. It never sees the hidden suite. The estimand throughout is **union recall at K** : the fraction of defective instances blocked by at least one of K independent readings, averaged exactly over all $\binom{K_{\max}}{K}$ subsets rather than by sampling. The matching quantity on the correct stratum is the union false-positive rate. A union recall figure is reported only beside the false-positive rate it cost.

Four auditor routes appear below. The **shipped cross-vendor** auditor is `gpt-5.6-terra`; the **same-vendor** auditor is `claude-haiku-4-5`, the generator’s own family and size; the **stronger same-vendor** auditor is `claude-sonnet-4.6`; and a third family, `gpt-6-astra` at high reasoning, is read at $K = 4$. The auditor’s constitution, its prompt builder and every route’s capability card are in the auditor harness, a separate repository, at the commit named in the provenance file, whose per-reading records carry the hash of every prompt actually sent.

The routes are **not sampled alike**, and this bounds every

cross-family union-of- K contrast in the paper. The cross-vendor route’s capability card does not permit a sampling temperature, so no sampling parameter is sent and the route varies between readings; the same-vendor routes permit one and the layer sends temperature 0, so those routes are near-deterministic. A union over readings that do not differ cannot grow, so the two families’ unions are not equally free to rise with K , and a cross-family difference in union recall confounds auditing ability with how much each union was permitted to grow. The asymmetry is a property of the shipped routing layer we measured rather than a design choice, we did not discover it until after the runs, and we report the $K = 1$ contrast beside the $K = 8$ contrast wherever it matters. A temperature-matched replication is the first item in §4.5.

2.3. Statistics

The unit of observation is the instance and the resampling unit is the problem, because problems recur across generation batches and instances sharing a problem are not independent. The K readings of an instance are repeated measurements of it, not K further independent observations. The bootstrap over problems estimates uncertainty from which tasks the benchmark happens to contain; it does not estimate how much a fresh batch of model calls on these same tasks would move the numbers, which would need replicate runs we did not budget for at every arm. Every rate carries a 95% percentile bootstrap over whole problem clusters (10,000 resamples, seeds fixed in each preregistration), with the Wilson interval printed beside it and marked as too narrow. Paired contrasts carry an exact McNemar test and a cluster-level sign-flip permutation test; where every discordant pair points one way the percentile bootstrap’s bound at zero is an artefact and an unconditional interval (grid-approximated, unclustered) is quoted instead — with one exception: the clarification study’s registration fixed the bootstrap for a contrast that turned out one-signed, so its interval is reported as registered and read beside its exact cluster test (§3.5). A saturation curve is summarised by a constrained exponential fit, and the fitted asymptote is quoted only when the preregistered flattening bar is met; otherwise it is labelled an extrapolation and the raw union at $K_{\max}$ is the number quoted. Counts of five or fewer are quoted as counts. Multiplicity is handled per study as each registered it: the study carrying the primary contrasts corrects over every comparison it computed, a Bonferroni family of sixteen with a threshold of 0.00313 that two comparisons clear; others registered a single primary, and study 19 reports its twenty further quantities without correction. **Interval coverage is measured rather than assumed or waved at.** Under the two strata’s actual cluster profiles — 110 instances in 56 problems, and 150 in 143 — with within-problem concordance 0.5, the percentile bootstrap’s cover-

age of a nominal 95% rate interval is, over 400 replicates of 600 resamples each (Monte Carlo standard error about 0.011 near 95% coverage and 0.020 at the lowest cell; single rates only, with integer-index percentiles where the analyses interpolate; `analysis/coverage_simulation.py`):

true rate	0.03	0.05	0.10	0.16	0.30	0.48	0.60
defect (56)	0.798	0.870	0.933	0.933	0.953	0.935	0.928
correct (143)	0.888	0.918	0.938	0.933	0.943	0.955	0.965

In these simulations coverage is near nominal at the true rates this paper’s headlines sit at and **badly anticonservative at small true rates**: at a true rate of 0.03 on the defect stratum the simulated interval covers about 0.80, not 0.95. This is evidence of potential undercoverage for analogous intervals, not a measured coverage of any interval quoted here. Intervals on rates below roughly 0.10 should be read as narrower than they claim. *An earlier version of this paper said the bootstrap under-covers “by roughly 2 to 5 points” and that its coverage under the clustered design was unvalidated. The first was a generalisation the source report had already withdrawn as unsupported; the second was wrong, since a clustered simulation sat in this paper’s own test suite. A pre-submission referee found both.* The simulation is `analysis/coverage_simulation.py` and costs no model calls. Two earlier paired-contrast constructions in this programme were found by review to cover 0.416 and 0.075 where 0.95 was claimed, and both were replaced. No sample size was justified in advance by a minimum-detectable-effect calculation.

An interval that contains zero is reported as one that does not exclude zero, and nothing more. No equivalence margin was registered anywhere in this programme and no equivalence test was performed, so none of those intervals is evidence that two conditions are alike or that an intervention did nothing — a distinction this work has already had to make the hard way, having once inferred equal cost from an interval containing zero and withdrawn it on review. Where a hypothesis fails its registered criterion we say it failed, not that the effect is absent.

All analysis is plain Python 3.13 using only the standard library — `random`, `math`, `itertools`, `csv` and `json`. The bootstrap, the exact McNemar test, the cluster sign-flip permutation and the exact subset enumeration behind every union-at- K figure are implemented in the repository rather than called from a numerical package, so each is readable at the point where it is used, and none depends on a third-party version.

2.4. Preregistration and independent review

Each study that made model calls was preregistered before its first one, fixing its population, hypotheses, primary

outcome, intervals, kill conditions, budget and stopping rule — **with two exceptions this paper reports rather than elides**. First, the severity sweep of §3.2 re-grades findings already archived, makes no model calls, and was never registered: every figure it produces is post hoc and exploratory, and none is a confirmatory test. An earlier version of this paragraph named only the second exception. Second, the analysis behind §3.5’s third-rating result was registered a day *after* its contrast had been computed and committed, in a document that claimed the opposite; that document is withdrawn in the repository and the analysis is registered nowhere. Its method was fixed in advance by an earlier amendment, and the disjoint reading, the abstention-excluded secondary and the broken-sheet comparison were not. Departures are numbered amendments committed before the step they govern, each recording what was known at the time; where an amendment followed an interim observation, the results say so and argue whether the change could favour the hypothesis. Each report was then reviewed by a model from a different vendor to the generator, with read access to the raw records, asked to re-run the analysis, recompute the intervals independently, and check every sentence against the records. **That reviewer is gpt-6-astra, which is also one of the auditor families measured in §3.2 and one of the two raters behind §3.4’s labels, and one of the two adjudicators whose joint labels define the clarification study’s outcome.** For most claims this is a cross-vendor check on work the reviewer had no part in. **For the claims that rank astra against the other families it is not: there the reviewer is the subject,** and those comparisons, and the clarification study’s adjudication, should be read as self-assessed until a model that audited nothing in this programme has checked them. This project’s own P3 Amendment 10 disqualified the same model as the clarification study’s manipulation-check rater for the same reason, yet kept it as that study’s second adjudicator and its reviewer; the standard was applied unevenly within one study, and not applied to the review at all until a pre-submission referee pointed it out. Reviews are archived verbatim, including those that refused approval, and several studies needed many rounds: on one day five reports were refused at once, and clearing those five took thirty-one further rounds. What moved, in most of them, was what the reports claimed about their measurements — and in most of the later refusals the object of the refusal was a repair written to satisfy the previous round, not the study it repaired. We state that because it is the most transferable thing this process produced: a fix written to satisfy a reviewer is itself an unreviewed claim.

The reviewer is a language model, and this paper is about the ceiling on language models reading code they did not write. That applies to our own instrument, and the round counts are the evidence: the studies here cleared review at

rounds **3, 4, 6, 7, 9, 10 and 11**, and in every case the round before the last still returned substantive findings. Some of those findings were defects introduced by the repairs between rounds, so the counts do not estimate a single pass's recall; they show that no report cleared in fewer than three. *Quotable at round N* means a report survived N passes, not that it was verified, and the marks in this paper should be read as a filter of unknown recall rather than as proof.

Two things sharpen that. First, the split. Measurements did move on review: the severity sweep's first review found its loader misreading one family's draws and another's population, and every cross-family figure was recomputed; the clarification study's second review moved its primary from 9 of 31 to 9 of 32 instances and its contrast from +29.0 to +28.1 points, by finding an instance wrongly excluded; and a label count inside a reviewed report was 71 where its data say 67. An earlier version of this sentence said measurements never moved, and a later one confined the exceptions to one study; both were wrong. Once a figure's inputs were right, the reviewers reproduced rates, intervals, subset enumerations and seeds independently and found no arithmetic error, and most refusals were about what the prose claimed. On recomputation the instrument is strong; on interpretation it is weak, and the two should not be quoted as one assurance. Second, the class of defect this leaves open is the one that looks right to a language model. The clarification study audited the wrong program for 384 readings — a per-problem field used where a per-instance one was needed — and every automated gate passed, because each checked that the code was *identical across conditions* and it was, identically wrong. It was found by hand, reading an adjudication sheet and noticing that a displayed program returned exactly what its hidden suite was said to reject. Neither the gates nor a reviewer found it, and we do not know what else of that kind remains.

Every review prompt in this programme stated the claimed result before the reviewer began. We attempted to test what that costs, and the attempt failed in a way worth reporting. One report, quotable at round 6 after five refusals and unchanged since, was dispatched again under a prompt whose own text withheld every figure and every prior verdict and asked for an independent computation first. **That prompt's list of binding artefacts then named two files that state those figures**, one of them the claims-ledger entry for the result under review. The blinding was not enforced, so the agreement between that pass and the briefed one establishes nothing, and we draw no inference from it. What the pass did return is a *disagreement count*: six substantive findings that six briefed rounds on the same file had not produced, among them that the report's two groups were two different definitions of *caught* rather than one population with duplicate rows, that its rating sheet displays at most three failing inputs while 81 of its 121 entries have

more recorded, and that a count of changed labels was 71 where the data say 67. ****That count is $n = 1$, uncontrolled and unregistered****: a seventh round might have found new things under any prompt, and no matched briefed round was run against the same file. We report it as a reason to test blinding properly rather than as a result about blinding, and the first round of every review after this one is dispatched with the conclusions withheld from the prompt *and* from every artefact it names.

Every study drawn on here has completed that review. One further study was re-run twice and closed without a quotable number; the paper names it only to say what it no longer supports.

3. Results

The claims this paper defends, each with its interval and its label, are listed in `CLAIMS.md`. A claim enters that file only after its report has passed an independent review by a model from a different vendor with access to the raw records. Every section below draws only on reports that have passed such a review. Section 3.5 reports seven studies whose reviews took between three and eleven rounds each to reach that point, and states what each of them does and does not license; none of them is load-bearing for §3.1–§3.4. Rates carry the 95% problem-cluster percentile bootstrap as the primary interval, with the Wilson interval beside it where the source report gives one, because problems recur across generation batches and an interval that treats instances as independent is too narrow. For single rates, simulations of an analogous interval under this clustered design (§2.3) suggest coverage near nominal at the headlines' rates and possible undercoverage below a true rate of about 0.10, under an assumed within-problem correlation; they do not measure the coverage of the intervals reported here, and the paired contrasts and category-conditioned shares have no measured coverage.

Sixteen comparisons were computed in the study that carries the primary contrasts, under one Bonferroni family with a threshold of 0.00313. **That correction is applied to p values and never to intervals**: every interval in this paper is an unadjusted 95% interval, so a directed statement resting on an interval alone, with no p value that clears its family, is worded as a direction and not as a finding. **Each p value belongs to the family of the study that computed it**. Those in §3.3 are members of that sixteen-member family. The one in §3.2 is study 18's single registered primary, which had no family to correct within; it would also clear 0.00313. The clarification study's are reported beside a primary registered on its interval, and no correction was registered for them. Where a p value appears below, the exact McNemar value and the cluster-level sign-flip permutation value are both given, because instances within a problem are not

independent and only the second respects that.

Table 1 collects every primary quantity below with its population and its label; the prose gives each one where it is used.

3.1. Returns to repeated reading diminish short of complete

The shipped cross-vendor auditor reads an increment holistically and returns findings; an increment is blocked if any finding is graded BLOCKER. Over eight independent readings of the same 110 defective instances, union recall rises from 10.7% at one reading to 30.0% [20.0, 40.7] at eight (Figure 1). The eighth reading still adds 1.93 points, cluster [1.14, 2.78], against the flattening bar of 1.0 points registered before the first model call. The bar was not met. Under exact subset averaging that statistic has a closed form: it is the number of instances flagged by exactly one of the K readings divided by Kn , here $17/(8 \cdot 110)$, so it is a count of singletons and not a fitted slope. Its interval resamples problems and so carries task-sampling variance; the statistic is made of reading-sampling variance, which the interval does not describe, and the bar should be read as a statement about these 17 instances rather than about a population slope. The constrained fit puts the asymptote at 31.5%, but on our own registered criterion that figure is an extrapolation and the raw union at eight readings is the number to quote. What the data support is diminishing returns at the budget we could afford, not saturation: at 1.9 points per reading the curve was still climbing when we stopped. That bar is also a function of where a family was stopped: exact subset-averaged increments do not grow with K — this curve's is 2.9 points at the fourth reading and 1.9 at the eighth — so a fixed bar is harder to meet at $K = 4$ than at $K = 8$, and meeting it at either says little by itself about where the curve ends. Twenty readings pooled across three auditor families reach 48.2% [36.7, 60.0], and the same twenty readings flag 54 of 150 correct increments, **36.0%** [28.2, 44.2] — more than twice the shipped route's 16.0%, so pooling buys its coverage at its own price. **Those twenty readings are not twenty opportunities.** The pool is eight cross-vendor, eight same-vendor and four *astra* readings, and the same-vendor route gains **1.8 points** from its first reading to its eighth against the cross-vendor route's 19.3; it is sent temperature 0 and returns nearly the same verdict each time, though how much of that is the temperature and how much the model this design cannot say. Counting that route as roughly one reading puts the depth nearer thirteen than twenty — a heuristic, not an estimate — and the pooled figure should be read as a union over three routes rather than over twenty independent looks.

The price is measured on the same footing. Union false positives on 150 correct increments reach 16.0% [10.1, 22.3]

at eight readings. The recall bought per point of false positive — recall gained over one reading divided by false positives gained over one reading — is **1.68** at eight, and it does not worsen as readings accumulate: the ladder is 1.67, 1.63, 1.64, 1.65, 1.66, 1.67, 1.68 from $K = 2$ to $K = 8$. An earlier version of this paper said it falls monotonically; the source table contradicts that — the ratio falls once and then rises, so repeated reading on this route buys recall at a slightly improving price, not a worsening one. It is a ratio of two measured rates and carries no interval of its own.

3.2. Which model reads, and under which severity rule

Replacing the auditor with a stronger model of the generator's own vendor moved union recall at eight readings by -26.4 points, cluster $[-37.3, -15.6]$: 4 of 110 against 33 of 110. Three instances were flagged only by the new family and thirty-two only by the shipped auditor, at an exact McNemar p of 4.2×10^{-7} and a cluster sign-flip p of 2.0×10^{-5} . This was study 18's single preregistered primary, so no family correction applies. Three things qualify what it means, all of them measured here and none of them favourable to reading it as a statement about auditing ability.

First, the two auditors are not at the same operating point, which is the comparison this paper elsewhere warns against. The stronger model flags correct code at 3.3% — 5 of 150, a count the statistics section says to quote as a count — whose bootstrap interval [0.7, 6.7] sits in the range where simulations of an analogous interval cover about 0.89 at a true rate of 0.03 rather than 0.95, so it may be narrower than it claims. Against the shipped auditor's 16.0% [10.1, 22.3]. It blocks less of everything.

Second, the sign reverses under a weaker severity rule (Figure 2). Counting any finding rather than only findings graded BLOCKER — an exploratory rule, not preregistered — the stronger model reaches 59.1% [47.3, 70.6] on the defect population at 34.0% [26.5, 41.7] on correct code, against the shipped auditor's 34.5% [23.6, 45.5] at 19.3% [12.8, 26.3]. The ratio of those two rates — a derived point ratio we carry no interval for — is 1.74 and 1.79 respectively, so nothing here shows one family buying more recall per false positive than the other. Under any-finding grading the stronger model's findings cover more instances and under the BLOCKER rule fewer, so the -26.4 depends on the counting rule. That does not show why the families' findings differ, or that any finding identifies the hidden failure.

Third, the two families are not sampled alike. The shipped cross-vendor route's capability card does not permit a temperature, so no sampling parameter is sent and the route varies between readings; the same-vendor route permits one and the layer sends temperature 0, so that route is near-deterministic. A union over near-deterministic readings cannot grow, and the two unions are therefore not equally

Table 1. Every primary quantity this paper quotes, with the population it is measured on and the label its source report gives it. Values are percentages or percentage points. Intervals are 95% problem-cluster percentile bootstraps unless the label names another method. The coverage simulation of §2.3 measures single-rate intervals only; the contrasts and the category-conditioned shares have no measured coverage. Every row is read from a committed record by `figures/make_table1.py`.

quantity	population	value	interval	label
Union recall, shipped route, $K = 8$	110 defect, 56 problems	30.0%	[20.0, 40.7]	preregistered
Union false positives, shipped, $K = 8$	150 correct, 143 problems	16.0%	[10.1, 22.3]	preregistered
Pooled union, 3 routes	110 defect	48.2%	[36.7, 60.0]	preregistered
its false positives	150 correct	36.0%	[28.2, 44.2]	derived (pending)
Gain from the eighth reading	110 defect	+1.93	[1.14, 2.78]	bar 1.0, not met
Same-vendor stronger — cross, recall	110 defect	−26.4	[−37.3, −15.6]	preregistered
its false-positive difference	150 correct	−12.7	[−19.3, −6.1]	prereg.; 3.3% v 16.0%
Rulebook rule — shipped, recall	110 defect	+30.9	[+19.1, +43.1]	preregistered
its false-positive cost	150 correct	+20.7	[+12.8, +28.8]	preregistered
Rulebook rule, flags	56 loop defect	+26.8	[+13.6, +40.4]	exploratory
its cost on correct code	56 loop correct	+12.5	[−3.6, +26.8]	unconditional, approx.
isolating contrast, pass rate	112 loop, 96 problems	+5.4	[−0.9, +12.1]	registered, failed
astra — cross, common $K = 4$	110 defect, 56 problems	+11.5	[+2.9, +21.3]	exploratory
its false-positive difference	150 correct	−0.3	[−3.8, +3.4]	exploratory
Undetermined, sheet groups	121 entries, 110 instances	+19.8	[+1.2, +38.1]	registered method
one definition of caught	57 against 53	+23.3	[+2.0, +43.6]	post hoc
Undetermined share, 3-family residual	44 of 57	77.2%	[62.1, 91.1]	post hoc
6-family residual	25 of 32	78.1%	—	post hoc (C14)
Clarified — original, diagnosis	32 selected, 19 problems	+28.1	[+7.7, +51.4]	one-signed; exact $p .0625$

free to rise. They do not: over $K = 1$ to 8 the cross-vendor family gains 19.3 points and the same-vendor family gains 3.0. At one reading the gap between them is −10.1 points, not −26.4; the rest accumulates as the two unions grow at different rates, and because model and sampling differ together, the design cannot say how much of that difference in growth is the sampling and how much the model. We regard this as the most serious limitation of the contrast, and §4.5 records what would be needed to remove it.

A third family argues the other way and belongs beside the headline. A high-reasoning model of the cross-vendor family reaches 32.7% [20.7, 45.0] union recall at four readings while flagging 10.7% [5.9, 16.1] of correct code, against the shipped auditor’s 30.0% at 16.0% with eight. Paired on the same instances, neither advantage is established together with the other: at a common four readings it leads on recall, +11.5 [+2.9, +21.3], with an unresolved false-positive difference of −0.3 [−3.8, +3.4], and against the shipped auditor’s full eight it leads on false positives, −5.3 [−10.3, −0.7] unadjusted, but not on recall, +2.7 [−6.5, +13.0] (§3.5). So the evidence does not support “a stronger model is a worse auditor”. It supports the narrower statement that this particular same-vendor model, at this severity calibration and sampled at temperature 0, blocks far less.

A same-vendor route of the generator’s own size reached lower union recall than the cross-vendor one on these instances, −12.7 points [−25.0, −0.9] at eight readings, signed so that a negative value favours the cross-vendor route. The interval is unadjusted and its upper end is 0.9

points from zero, and that route is sent temperature 0, so this is a direction under this sampling and not a ranking of auditors. We reported until 2026-09-21 that a study on a second substrate reverses that sign. *That is withdrawn*: the study closed without a quotable number, and on its clean run the same contrast included zero, so no reversal was established either (§3.5). The figure above is therefore a measurement on this substrate, and whether it is a property of same-vendor auditing is open — which is what it was before we attempted the second substrate, not something that study settled in either direction.

3.3. One rulebook sentence moves flags, and does not measurably move outcomes

Adding a single rule to the auditor’s constitution, telling it to find what the visible tests do not cover, moved flags on the defect population from 10 of 56 to 25 of 56: +26.8 points, cluster [13.6, 40.4], exact McNemar $p = 0.0003$, cluster sign-flip $p = 0.0009$. With the pooled defect-and-correct flag contrast this is one of only two comparisons that clear the Bonferroni threshold, and the two are the same effect seen twice.

The estimand is narrower than the sections above and must be read at its own scale. It is measured on 56 instances from 41 problems, not the 110-instance defect population; it counts flags, not union recall against the hidden suite; it is one reading, because every loop arm revises once, not a union over eight; and it is exploratory, because the preregistration named outcome contrasts rather than flag contrasts. For those reasons we do not compare its magnitude with




Figure 1. Returns to repeated reading diminish short of complete. Union rate against the number of independent readings K , for two auditor families. (a) recall on defective increments; (b) false positives on correct increments. Bars at the right of each panel give the 95% problem-cluster bootstrap interval of each series at $K = 8$, dodged in x so overlapping intervals stay separable. The cross-vendor curve does not meet the preregistered flattening bar — its last-step gain is 1.93 points [1.14, 2.78] against a bar of 1.0 — and is not described as saturating. The same-vendor curve meets it, at 0.23 points, but that route is sent temperature 0 and its verdict varies across readings on only 3 of 110 defective instances under this configuration, so its flatness is not evidence of exhausted headroom; whether it reflects the sampling or the model is not separated here (§3.2). An earlier version of this caption said neither curve met the bar, and the figure carried an annotation left over from the withdrawn second substrate. A second substrate was plotted here until 2026-09-15; it was re-run twice and closed without a quotable number (§3.5).

the -26.4 or the $+19.3$ of the preceding sections, which are union recall at $K = 8$ over 110 instances. They are different quantities on different populations.

Two costs travel with it, and the first is weaker than an earlier version of this paper said. On correct code the same rule moved flags by $+12.50$ points. ****Every discordant pair points the same way**** (7 against 0), so by the rule stated in §2.3 the percentile bootstrap’s bound at zero is an artefact and the admissible interval is the unconditional one — computed with a grid-approximated nuisance supremum and ignoring clustering, so approximately rather than guaranteed exact — which is $[-3.6, +26.8]$ and **does not exclude zero**; its exact McNemar and cluster sign-flip p are both 0.0156, against that study’s Bonferroni threshold of 0.003125. Earlier versions of this paper quoted the forbidden bootstrap interval [5.17, 21.82] in four places, including the abstract. The directed statement the records support is the pooled one, where 23 discordant pairs point one way and 1 the other, exact McNemar $p = 3.0 \times 10^{-6}$ and cluster sign-flip $p = 1.0 \times 10^{-5}$: the sentence makes the auditor louder overall. On the correct stratum alone the cost is $+12.50$ points with an interval that contains zero. And on the contrast that actually isolates the rule — the referent arm against the cross-vendor arm on hidden-suite pass rate after one revision, over both strata, 112 instances on 96 problems — the effect is $+5.36$ points, cluster $[-0.89, +12.07]$, exact

McNemar $p = 0.146$ and cluster sign-flip $p = 0.183$, which does not exclude zero. The rule demonstrably moves what the auditor writes down. Whether that converts into code that passes the hidden suite is not established, and since this paper argues elsewhere that a flag is not a catch, that distinction applies to our own intervention first.

3.4. Most of what the auditor missed is failure the specification never determined

Fifty-seven of the 110 defective instances were flagged by no reading of the three families read at that point. That is a statement about twenty readings, not about detectability: if an instance is caught with probability p on each independent reading, twenty readings miss it with probability $(1 - p)^{20}$, which is still 98% at $p = 0.001$ and 36% at $p = 0.05$. Never flagged here therefore does not mean never flaggable, and we use the residual as a population to characterise rather than as a set of defects shown to be invisible.

It is not never flaggable, and we can put a number on that. Reading six families at forty readings, including the same auditor under a rulebook with one sentence added, shrinks the never-flagged set from 57 to **32 of 110**, and **43 of the 68 instances both raters call specification-undetermined are flagged by some reading of some family**. Being undetermined is not sufficient for being missed.




Figure 2. Union recall is a property of the decision rule, not only of the auditor. Each point is one auditor family re-graded under one of four decision rules on the same archived readings, with problem-cluster 95% intervals on both axes; no new model calls were made. **The families are read at different depths:** *cross*, *self* and *self-strong* at every one of their eight readings, *astra* and *self-frontier* at their four, so the figure places each family at its full archived depth and is not an equal-depth comparison; the common-depth contrasts are in Table 1. The dotted guides join a family’s four points and are *not* a measured curve — only the four points exist. Two things the accompanying table cannot show: the families do not lie on one frontier, and *self-strong* has no reachable point between 3.3% and 34.0% false positives, which is exactly where a comparison matched to the cross-vendor auditor would sit, so its ranking there is bracketed rather than measured. Nothing in this figure assesses finding content; a flag is not a diagnosis.

What does not move is the composition: the undetermined share of the never-flagged set is **44 of 57**, 77.2%, at three families and twenty readings, and **25 of 32**, 78.1%, at six families and forty readings. The residual loses nearly half its members and keeps its shape. *An earlier version of this section reported only the 57 and described it as flagged by no reading of any family, which was true of three families and false of the six this programme measured; a pre-submission referee found it.*

The first classification of that residual asked whether the visible tests exercised the failing input before it asked anything else. It called 46 of 57 unexercised edges, and an earlier draft of this paper rested a claim on that.

Re-rating the same instances under a rule that asks first whether the specification’s prose determines the hidden suite’s expected value reverses it. Two raters agree that 44 of 57 (77.2%, cluster [62.1, 91.1]) are cases where the prose does not settle the value, and that 2 of 57 are cases where

it does and the visible tests merely never built the input. The preregistered condition for withdrawing “the residual is mostly unexercised edge” fires on both of its clauses, and the claim is withdrawn. The interval is conditional on a category and has no committed coverage simulation, so it is uncalibrated, as is every category-conditioned interval we report.

A classification that reverses when the order of its questions changes is, by itself, weak evidence about the auditor, and the raters are not independent of it: one was the author, who wrote the rubric and knew the earlier labels and which claim was at stake. Three things support the second reading. The rubric was fixed before any label was written. The labels were committed before we looked for outside evidence. And when a first pass was found to have carried instance identifiers and a sheet name into the raters’ view, the labelling was re-run with identities stripped and the result held. Where an outside party had independently recorded that one of these specifications is defective, in twelve MBPP tasks named in a paper written for another purpose, our raters agree on 9 of the 11 that fall inside our defect population, with both disagreements on the boundary that paper’s own “incomplete” category draws. That check is one-sided by construction: it can support agreement where an outsider looked, and can bound nothing about the rate.

A rater outside the residual’s families, asked a different question. One of the two raters above is *gpt-6-astra*, whose own misses help define the residual. A model that is not one of the families defining it, *gpt-5.6-luna* — which read these instances only as a secondary audit route in an earlier architecture study — later rated the same 57 instances, post hoc, under a different rubric — whether the prose determines an answer at all, on a sheet showing at most three failing inputs and never the expected values (§3.5). It labels **38 of 57 undetermined**, 13 determined and 6 cannot-tell. The two figures measure different constructs: the earlier category also admits prose that determines an answer the oracle contradicts, which this rubric does not count as undetermined, so 38 is not a corrected estimate of 44. What they share is that a participating and a non-participating rater both put a majority of the residual on the undetermined side, and the second agreed with itself on only 7 of 11 repeated items.

3.5. Seven studies, and what they license after between three and eleven review rounds each

Four further studies had reported when this paper was first written, and all four were refused quotation by their independent reviews on the same day. They have since passed, after between four and eleven rounds each. Three more have since reported and passed, at seven rounds, six and three. We report them here rather than folding them into the sec-




Figure 3. The residual contrast, printed beside the instrument that produced it. a, The share of instances a rater judged specification-undetermined, missed group minus caught group, under each reading, with problem-cluster 95% intervals. The two sheet groups use different definitions of *caught* and share eleven instances, so the two readings that apply one definition are the ones that compare populations; both are *post hoc*. The red point is the same rater on the same sheet before its control arm was rebuilt, when 42 of 53 caught entries carried no failing-input evidence and all 42 drew *cannot-tell*: an unanswerable control arm did not produce noise, it produced a larger contrast pointing the way the hypothesis points, and a preregistered abstention gate is what excluded it. **b**, The same rater’s agreement with itself on the eleven instances the sheet carries twice, with identical text and identical evidence. **The contrast in a is read off the instrument in b**, and the claim this paper draws from it is stated under that constraint.

tions above because three of the four are narrow or negative results whose limits are part of the claim, and because the shape of the refusals is itself a finding about this method.

A rulebook sentence moves a second vendor’s family. The referent rule — *find what the visible tests do not exercise* — raises union BLOCKER recall on *self-strong* (Sonnet 4.6) at four readings by **+20.9 points**, cluster [+10.8, +31.8], at a false-positive cost on correct code of **+8.0 points** [+3.4, +12.7]. The grading rule shows no demonstrated improvement: +0.9 points [−1.8, +4.5], and that hypothesis fails its registered criterion, which demonstrates neither equivalence nor that the rule did nothing. This reproduces §3.3’s effect on a different vendor’s model at a different *K*, with its own price. A post-hoc quantity from the same study — whether a flagged finding *reports a failure* on the named class — is **3 of 110** under the shipped constitution and **27 of 110** under the rule, and *may not be quoted as defect recognition*: its rubric admits a finding even where it grades the class non-blocking or says the specification is silent, so findings that disclaim the requirement are counted in. Seven review rounds passed over the opposite error, in which the report called that instrument *narrower* than defect assertion — the direction that makes a quantity look more conservative than it is.

A measured 60.9% is not a higher ceiling, and it is the same estimator as the 30.0%. On the defect population the same rule takes union recall — the identical quantity, blocked by at least one BLOCKER, unioned over eight readings — from 33 of 110 to **67 of 110 = 60.9%** [48.6, 72.5], **+30.9 points** [19.1, 43.1], at **+20.7 points** [12.8, 28.8] of

false positives, with a single-draw rate of 23.6% [17.1, 30.3] that fails the product bar on all eight draws. *Earlier versions of this paper called 33 of 110 union recall and 67 of 110 flag coverage, and a pre-submission referee observed that the name tracked which number suited the framing.* One name is used here for one estimator, and The caution that a flag is not a catch attaches to both equally. What separates them is not their status as measurements but their price: 36.7% [28.5, 44.8] false positives against 16.0%, which is why the shipped route is the shipped route. The fitted asymptotes differ by +27.1 points [10.9, 39.0] and *that does not establish a higher ceiling*: the cross-vendor arm never flattened, so one of the two terms is an extrapolation; heterogeneous low-probability detection imitates a ceiling at this budget, since an instance found with probability 0.02 per reading survives eight readings about 85% of the time; and this study’s own zero-inflated estimator finds no never-detectable class in *either* arm. Its adjudication kill is arithmetic rather than a labelling artefact — only 9 of 110 instances drew any finding, so the largest defect-asserting count available under perfect adjudication was 9, below the floor of 20. Every adjudication-derived figure in it is conditional on labels whose production can no longer be audited: one rater’s raw reply, prompt, launcher and execution log lived only in a session scratchpad and were destroyed between the review that flagged them and the repair.

Corroboration does not validate a generated test. A wrong generated test — one that fails on the dataset’s canonical solution — is not reliably recognised by asking whether further independent generation draws also fail on the same

candidate. Neither paid rule reaches the preregistered threshold on the registered population — failing applications on all three candidate strata, 58 of the 107 on candidates that already fail their visible tests: retained wrong-test rates are **12.2%** (Wilson 7.0–20.6; cluster 3.6–22.6) and **14.9%** against a limit of **2%**, and neither is shown to improve on the free within-draw comparator at 12.8%: the better paid rule keeps the same wrong applications plus four correct ones, a small observed advantage the study cannot resolve. On the candidates a product would meet — those passing their visible tests, post hoc and imprecise — the rates are 18.2% and 23.3%. A floor of 2.4% holds on the original unique-test quantity because 29 of the 45 wrong tests never fail on any candidate at all. The claim is negative and substrate-bound: these rules do not validate generated tests here, not that generated tests cannot be validated.

An injected population does not give the residual result a prospective test. We constructed a population of injected defects intended to make specification-determinedness hold by construction. It does not. The filter registered to recover each edit’s first failing hidden input is a truthiness check on a field the injector wrote; of 281 recorded witnesses 179 are strings and 102 are lists, none is executed, and four of the remaining five filters accept cases their registered wording excludes — a suite that reports success after an early exit, a process that terminates without an assertion or an exception, an added import from a module already imported. *No filter establishes specification entailment, and repairing the first one alone would not have:* recovering the failing input and executing a witness shows the edit changes behaviour on an input the hidden suite exercises, not that the specification entails the value expected there. What may be quoted is descriptive. At eight readings the auditor blocked **90 of 92** accepted injected instances against **11 of their 92** unmodified twins, 79 discordant pairs one way and none the other, and against **33 of 110** on the natural residual; the instances the construction’s own gate *refused* were blocked **31 of 40**, which is why that gate cannot be assumed conservative. A probe separates the injected set from the natural one at 96.7% on the items it answered, *by a feature it does not identify*: the two arms share no problem, so task composition alone could produce the result, and the probe’s model is one of the two gate models. So §3.4 keeps exactly the status its own study gave it, post hoc and unreplicated. A successor needs an independent party to validate the actual violation *and* the specification-entailing expected behaviour under a fixed, blinded protocol; a second substrate alone does not suffice, and an independent rater alone does not either, because task composition and construction confounding survive both.

A fifth study was re-run twice and closed without a quotable number. A second benchmark was reported here until its review refused quotation and we confirmed the rea-

son: the visible-test text shown to the auditor *and* to the generator was sliced out of its test class without the class header, so it failed to parse on all 300 tasks while scoring executed the intact class. The repair for that fault introduced a second one, re-attaching a hidden method’s decorators to the first visible method on 26 tasks. A third run on a clean extraction reproduced five of the six findings that the second run had reversed and did not reproduce the sixth; the criterion registered before those numbers existed was stated over the set, so the strict reading applies and the study is closed. Three runs, about \$82, and no quotable number. Whether five of six should have licensed the five is recorded as an open question in `CLAIMS.md` rather than settled by us.

Measured union recall is rule-dependent, and a headline’s magnitude belongs to its own rule. Re-grading findings already archived by two of the studies above under four decision rules, with no new model calls, leaves the shipped BLOCKER threshold’s estimate intact for its own rule — which *was* preregistered, as “ ≥ 1 BLOCKER finding” — while *reversing its sign* under any-finding grading. That reversal is a mechanical consequence of changing what counts, and it identifies nothing about why the families produced different findings or different severities. At a common depth of four readings, *astra* exceeds the cross-vendor auditor on union recall over the defect stratum by **+11.5 points** [+2.9, +21.3], paired within instance. **The two sides are not averaged alike:** the cross-vendor arm has eight readings and is averaged over all $\binom{8}{4} = 70$ four-draw subsets, while *astra* has exactly four, so $\binom{4}{4} = 1$ and its estimate is a single realisation with none of its reading-selection variance averaged out. An earlier version of this paper said 70 subsets on each side; the source report has it right and the manuscript lost it. The contrast should be read knowing that one arm’s luck in four draws is fully in the number and the other’s is not, with a false-positive difference of -0.3 points $[-3.8, +3.4]$. On the cross-vendor auditor’s complete eight-draw ladder the same contrast is $+2.7$ points $[-6.5, +13.0]$ with a false-positive difference of -5.3 points $[-10.3, -0.7]$: that comparison points toward *lower false positives, not higher recall*, on an interval whose upper end is 0.7 points from zero and which is not adjusted for the other contrasts computed beside it, so it is a direction the data lean and not one they establish. Observed false-positive rates near one another are not a matched operating point, and one family has no reachable point between 3.3% and 34.0% false positives — precisely where a matched comparison would sit — so its ranking is *bracketed rather than measured*. Nothing here assesses finding content against the hidden failure, so none of it speaks to what any finding identified.

The ambiguity gap survives a rater who is not an author, and the instrument that shows it disagrees with itself.

The contrast in §3.4 between missed and caught instances had been rated by a model and by an author. A third model re-rated the whole sheet after we rebuilt it on mechanically recovered witnesses — necessary because its control arm had been *unanswerable*: 42 of 53 caught entries carried no failing-input evidence and all 42 drew `cannot-tell`. The share labelled *undetermined* is 43 of 68 against 23 of 53, **+19.8 points** [+1.2, +38.1]; those two groups are not two populations, because they define *caught* differently and 11 instances fall in both, and removing those 11 from the missed group gives +23.3 points [+2.0, +43.6] — *post hoc*, and larger, so the overlap had understated it. The limit is part of the result: on those 11 instances, shown with identical text and identical evidence, **the rater agreed with itself on 7 of 11** on this binary outcome and 5 of 11 on the full label, where the earlier study’s two raters each agreed 11 of 11 (Figure 3). Every interval’s lower end lies within five points of zero, and this remains a model rather than the outside human the design asks for. Separately, the *unrepaired* sheet scored +54.2 points [+38.5, +69.5] — 2.7 times the repaired contrast, in the direction that flatters the hypothesis — which is evidence about sheets and not about specifications, and which a preregistered inconclusiveness gate catches.

Adding the rule the specification never settled moves diagnosis, where the rule is the oracle’s. The contrast above conditions on the specification; this study changes it. The population is the 44 never-flagged instances that both earlier raters called specification-undetermined; for the 32 on 19 problems that passed the generation gates, a model shown the failing *inputs*, but never their expected values, wrote one behavioural rule into the prose, and an auditor re-read the same code and the same visible tests under the original, the clarified and a same-length placebo specification, four readings each. An instance counts as diagnosed only if both adjudicators judged some finding to state the behaviour the hidden suite expects. Diagnosis rose from **0 of 32 to 9 of 32** under clarification and stayed at 0 of 32 under placebo, **+28.1 points**, problem-cluster 95% interval [+7.7, +51.4], which meets the registered criterion; the registered secondary, on the 23 instances where study 21’s consensus and the third rater’s label agree, is 8 of 23 against 0 of 23, +34.8 [+10.0, +60.9]. **Given the observed data, that criterion added little.** Every discordant pair points one way, so no bootstrap resample can fall below zero, and with successes on five of 19 problems a resample is exactly zero only when it misses all five, with probability $(14/19)^{19} \approx 0.003$; the registered run has 27 such resamples in 10,000. Once five success-bearing problems are observed, the interval’s exclusion of zero adds little to them — with one such problem it would not have excluded zero — and §2.3 says to read a one-signed contrast from an exact test instead. The registration kept the bootstrap and we report it as registered,

but the informative figure is the cluster sign-flip test the registration prefers: the nine sit on five problems, and it gives $p = 0.060$ (exactly $2/2^5 = 0.0625$ on enumeration), against 0.0039 for a McNemar test that counts nine independent pieces of evidence where five exist. It is not significant at 0.05. The registered naming-rate secondary was never computed. Two further limits are part of the result. *The inferred rule contradicts the hidden suite entirely on 23 of the 32 instances and partly on 2*, and diagnosis was 0 of 23 where it contradicted against 9 of 9 elsewhere. That split was classified after the outcome was seen, not assigned, and the agreeing instances may simply be the easier problems; it does not say which property of the edits moved diagnosis. And a separate model rated the clarified prose more determining than the original by +28.1 points [+2.6, +54.3] but than the placebo by +18.8 [−9.4, +45.2], so whether the placebo was inert as a specification is unresolved here. One adjudicator is an author and the other is `gpt-6-astra`, the model that also reviewed this study (§2.4); they agreed on 84 of 91 items, $\kappa = 0.85$, which is agreement and not accuracy. The study supports a specification-edit contrast on these selected instances. It does not establish that underdetermination caused the original misses, and §3.4 is not strengthened into that claim by it.

What the refusals were about. These seven studies cleared review at rounds 3, 4, 6, 7, 9, 10 and 11, fifty rounds in all; an earlier version said forty-four, which counted four of the studies only from their first refusal. Most of those rounds ended in a refusal of the *repair* rather than of the measurement. Measurements moved on review in three of the seven: the sweep’s first review found its loader reading one family’s draws 4 to 8 as 1 to 5 and another’s population from a partial draw, and every cross-family figure was re-computed; a label count in the third rating’s report was 71 where its data say 67; and the clarification study’s primary moved by one instance and three classification rows (§2.4). Every figure above reproduced independently at every round after its last repair. The recurring failure was a check or a sentence that claimed more than it delivered: a binding that protected one sentence of a withdrawal spread across six; the same binding rewritten to protect all six, with two of them truncated before the words that carried their meaning; a check that stripped HTML comments and called the result what a reader sees; and a test named for a code path that never called it. A second pattern appeared in the last thirteen rounds and is worth separating from the first: most of those refusals were *overstatements away from our own interest* — a registration whose chronology its own repository contradicted, a withdrawal that overcorrected past what the record supported, a publication prohibition a registered gate never imposed, a limitation credited to another study’s work that the other study had not performed. Modesty is not the safe direction; it is another direction in which a sentence can be

wrong. The corrective is not a better patch but a smaller claim, written where the claim was — and a heading, a title and a limitation are claims.

3.6. What follows for a recall number

Taken together, the results above say an audit recall figure is a joint property of the auditor and its material. It moves with which route reads — model and sampling together, which this design does not separate — with what the rulebook tells it to look for, and with the severity threshold at which a finding becomes a block, the last of which is rarely reported. What it misses is associated with whether the hidden suite asks a question the specification answers: observed in the residual, and in a specification-edit contrast on selected instances, not isolated as a cause.

4. Discussion

4.1. What a union-recall number means

The caution this subsection argued — that a union-recall figure is contingent on the rate it costs, so recall quoted without its false-positive rate is not portable — rested on a second substrate that was re-run twice and closed without a quotable number (§3.5). We state the caution as a methodological point rather than as a measured result, because the argument for it does not need our second substrate: recall and false positives move together with any severity threshold, and a critic can be tuned to any recall by blocking more. What our records still support directly is the narrower version of it on the substrate we report, where the same auditor at two severity rules sits at two different operating points (§3.2), and where the shipped auditor’s 30.0% [20.0, 40.7] recall is quoted only beside the 16.0% [10.1, 22.3] it cost.

The cleaner experiment for the general claim is a severity sweep within one substrate. Our sweep re-grades archived findings at four decision rules and shows each family’s recall and false-positive rate moving with the rule (Figure 2); because it re-grades rather than re-asks, it cannot place the families at a common operating point. Comparing two substrates was never that experiment: they differ in specification length, solution length, problem population, defect distribution and suite construction at once.

4.2. Repeated reading buys nothing where the verdict does not vary

Union-of- K is the instrument this paper uses to find a ceiling, and it is worth saying plainly what it cannot measure. Part of that needs no experiment at all: a union over readings that do not differ cannot grow, so a union-of- K curve is flat for a route that returns the same verdict however often it is asked. That is arithmetic, not a finding.

The empirical illustration we had for it came from the second substrate, which closed without a quotable number (§3.5), so we no longer offer one. What survives is the design fact behind it, which our own records do establish: the same-vendor routes are sent temperature 0 while the cross-vendor route is sent no sampling parameter at all (§3.2), and over $K = 1$ to 8 the first family’s union gains 3.0 points against the second’s 19.3. How much of that difference is sampling and how much is the operating point, we cannot say, because no temperature-matched arm was run. We report it as a confound on our own contrasts rather than as a result about auditors.

What follows for best-of- K , union-of- K or panel aggregation over critics is only the arithmetic: a route whose verdict does not vary between readings gains nothing by being read again. We measured per-instance flag counts, not any distance to a decision boundary, and ran no panel aggregation, so we offer no law for how much a varying route gains. Nor is the diminishing shape itself evidence of a ceiling: exact subset averaging makes each reading’s marginal gain no larger than the last by construction, and what the curves measure is the size of those gains.

4.3. What the ceiling is made of

Repetition is the smallest of the contributions we could measure. Eight readings of one auditor buy 19.3 points of union recall, and twenty readings across three families leave 51.8% [40.0, 63.3] of the defect population untouched — the complement of the pooled 48.2% [36.7, 60.0], which those twenty readings buy at 36.0% false positives on correct code. Changing which model reads moves the same quantity — union recall at $K = 8$ over the same 110 instances — by 26.4 points, though §3.2 gives three reasons not to read that as auditing ability. Changing one sentence of the rulebook moves flags by 26.8 points, but on 56 instances at one reading, and we do not place that number on the same scale as the other two.

A fourth contribution is not the auditor’s at all, and we can only describe it post hoc. Under a rule that asks whether the specification’s prose determines the hidden suite’s expected value, 44 of the 57 never-flagged defects fail that test. We attempted a constructed population where the prose does determine the answer; review established that no filter in it enforces that property, and what it licenses is descriptive only (§3.5). The association between never being flagged and never being determined is empirical, not definitional — selecting the missed instances does not fix their labels, and against caught instances the third rater’s post-hoc contrast is +23.3 points [+2.0, +43.6] — but selection and confounding prevent a causal reading, so it is an observation about where the misses sit and not an estimate of how much the benchmark contributes. One intervention changes the

specification instead of conditioning on it (§3.5): adding a behavioural rule to the prose moved diagnosis on 9 of 32 such instances, all of them where the added rule agreed with the hidden suite and none of the 23 where it contradicted it. It is the first evidence here that manipulates the specification instead of conditioning on it, though only within a population selected for having been missed and called undetermined, and it cannot separate a specification that settles the question from one that settles it correctly.

The constructed population carries a limitation we measured rather than assumed, and a second one that review established. A frontier model separates its edits from natural ones at 96.7% on the items it answered, between 79.9% and 97.3% once the items it would not classify are counted at either extreme — *by a feature it does not identify*, since the two arms share no problem and the model is one of the two that gated the population. And the construction does not enforce specification-determinedness at all, so it licenses *no* bound on the measured ceiling. We had written that it did. It does not.

4.4. Flagging is not naming

An audit that blocks an increment for the wrong reason has still blocked it, and a recall number counts that as a catch. With the finding texts kept for one reading of the shipped auditor and adjudicated blind by two raters, the auditor returned a finding on 9 of 110 defective instances and asserted the actual failing behaviour on 5 to 7 of them, depending on how the two raters are combined. The raters agreed at $\kappa = 0.391$, which we report beside the rate it produces, because a moderate agreement is part of what the rate is worth. That study’s preregistered kill condition fired on this result, and §3.5 states its headline in the form the kill requires. The figure is a property of one reading and is not a naming rate for the eight-reading union; establishing that would take adjudicating eight readings’ texts, which were not archived.

4.5. Limits

The confound we would fix first is that auditor families are not sampled alike: one route is sent temperature 0 and the other no sampling parameter at all, which bounds how far one union can grow before any difference in ability is measured. Every cross-family union-of-K contrast here inherits it. Second, no auditor is compared with another at a matched false-positive rate; our sweep re-grades archived findings at four decision rules, which cannot move a family to a point its findings do not reach — one has none between 3.3% and 34.0% false positives — so the paper cannot say whether the families differ once brought to the same operating point. Third, the primary intervals are problem-cluster percentile bootstraps over 56 clusters; a simulation

of an analogous single-rate interval (600 resamples rather than 10,000, tested rates only) covers near nominally at the headlines’ rates and anticonservatively below about 0.10 (§2.3), and for the paired contrasts and conditional shares it is unmeasured, and no sample size was justified in advance by a detectable-effect calculation. Fourth, the substrate is short, self-contained Python functions with executable hidden tests, so nothing here measures auditing a change to a large program with no oracle, and both of its benchmarks are old enough that memorisation cannot be ruled out. Fifth, the generator is a single model, so the defect population is one model’s failure mode. Sixth, the residual rubric’s raters were a model and the author, who wrote the rubric and was not blind to the earlier labels; one external label set agrees with them across eleven overlapping tasks. Seventh, our attempt to build a population of specification-determined defects did not enforce the property, and its edits are separable from natural code at 96.7% [79.9, 97.3] by something the probe does not identify. Each of these is recorded in the study that carries it, together with the number it would change.

4.6. What we would do next

The severity sweep we ran re-grades archived findings under four rules and shows the auditor contrast is rule-dependent (§3.2); it cannot reach a matched operating point for every family, because re-grading only moves a cut through findings already produced, and one family has no reachable point between 3.3% and 34.0% false positives. Re-asking rather than re-grading — instructing each family at graded severity and reading afresh — is the experiment this paper most needs, since only that can place the families at a common false-positive rate and decide whether the auditor contrast survives there. Replicating the cross-family arms with the sampling parameters matched would remove the confound above, and extending the cross-vendor family past eight readings would turn a statement about our budget into a statement about a curve. We have run a first version of the experiment we most wanted — hold the code and the visible tests fixed and vary only the specification (§3.5) — and it left two gaps that decide what it means. Its clarifier inferred the missing rule from failing inputs and was wrong on most instances, so the rule’s correctness was measured rather than manipulated; writing the rule from the oracle for some instances and a deliberately wrong rule of matched length for others would separate a specification that settles the question from one that settles it correctly. And its manipulation check did not show the placebo inert; a placebo rated no more determining than the original would. Both are cheap, because the code, the tests and the adjudication rubric already exist. Beyond that: adjudicating more than one reading of the shipped auditor would let a naming rate be compared with a union flag rate at the same K; a second generator would show whether the residual’s character is

a property of the benchmark or of one model’s failures; a population of natural defects independently known to be specification-determined would separate the injected population’s artificiality from the property it was built to isolate; and a substrate whose ground truth is not a hidden test suite would test the assumption every ceiling here rests on, which is that a test suite asked the right question.

5. Related work

Critics that read code. Trained critic models write natural-language critiques of model-written code and outperform unaided humans at catching bugs (McAleese et al., 2024). That work evaluates on both deliberately inserted bugs and naturally occurring ones: on code containing naturally occurring model errors its critiques are preferred to human critiques in 63% of cases, and it surfaces errors in training data previously rated flawless. It also reports that critics hallucinate bugs, and that a human–model team hallucinates less than the model alone. So the natural-defect setting and the precision cost of a louder critic are both already its territory, and we do not claim either as new. And the reported quantity is a preference or a per-critique comparison rather than the union recall of repeated independent readings against an executed oracle. We should be plain that this work already occupies two of our four axes: it reports a precision–recall trade-off for critics, and it separates tampered from natural bugs. What we add on this axis is narrow — the union of repeated independent readings as an instrument for locating a limit, ground truth supplied by an interpreter rather than by preference, and the false-positive rate carried beside every recall figure. Aggregating several review passes is known to help: multi-review aggregation on 1,000 verified pull requests raises F1 by up to 43.67% (Zeng et al., 2025). The motivation is shared with our union-of- K , but the operation is not: there an additional model call synthesises the sampled reviews, while our union counts an instance whenever any reading flags it, and the two estimate different things.

Repeated sampling and coverage. That coverage — the fraction of problems solved by at least one of K samples — rises steeply and then flattens is an established regularity of repeated sampling, reported over four orders of magnitude of K and fitted by a power law (Brown et al., 2024). Our union-of- K is that quantity with a flag rule in place of a correctness check, so §3.1 re-measures a known shape rather than discovering one, and at $K \leq 8$ our budget is small beside theirs. The contribution is not the shape. It is that the curve is read here as a measurement of a ceiling, reported beside the false-positive rate it costs, and tested against a registered flattening bar that it fails — which is why we report diminishing returns rather than a limit.

Self-assessment. That models are poor judges of their own

output is established from two directions: self-correction without external feedback does not reliably improve reasoning and can degrade it (Huang et al., 2024), and models prefer their own generations when used as judges (Chen et al., 2025). Our same-vendor and self-audit arms are the code-audit case of that literature, and we read our null there as consistent with it rather than as new. What our arms add is the false-positive rate beside the recall. We also reported that the direction reverses on a second substrate; that claim is withdrawn with the study it came from (§3.5).

What the hidden suite decides. EvalPlus rebuilt HumanEval and MBPP with larger suites and reported defects in the original reference solutions (Liu et al., 2023); its code carries bespoke oracles for tasks where differential testing against the reference was not applicable. Richter & Papadakis (2026) show that models collapse onto a single incorrect reading of an underspecified task and name twelve MBPP tasks they identified by hand; eleven fall inside our defect population and our raters, who labelled them before that paper was fetched, agree on nine. From the test-generation side, the dominant error in generated tests is a wrong expected value rather than an invalid input (Xiong et al., 2023), and generated oracles tend to capture a program’s actual rather than its intended behaviour (Konstantinou et al., 2024). Our claim is the auditing counterpart: where the prose does not fix the value, a recall number computed against that suite measures the benchmark as much as the auditor.

Consensus as a substitute for ground truth. Several systems validate generated tests without a reference by agreement, through self-consistency, dual execution, or multi-agent discussion for JUnit suites (Taherkhani et al., 2026; Xu et al., 2025). Both measure what survives the filter against a ground-truth implementation — 84–91% of retained tests valid in one (Taherkhani et al., 2026), 0.89–0.93 oracle correctness in the other (Xu et al., 2025). Our measurement differs in unit and in bar: the share of wrong tests among retained *failing test–candidate applications*, against a preregistered 2% threshold, over a frozen set of generated tests and candidates that includes candidates already failing their visible tests. Corroboration across independent generations does not reach it (§3.5), and wrong tests recur by problem. That recurrence is compatible with the self-consistent errors every detection method struggles with (Tan et al., 2025) — the same wrong answer produced again — but our records hold indices, not assertions, so they do not show that the recurring errors are the same one.

Who audits whom. That an evaluator should not be the model under evaluation is supported by measurements of self-preference in LLM judges against verifiable ground truth (Chen et al., 2025). What we measure is not self-preference: a frozen generator’s code audited by routes

that differ in model, operating point and sampling cannot separate preference from performance, so we report the same-vendor contrast as a contrast between routes, on one substrate.

Preregistration. Preregistering an evaluation protocol is beginning to appear in machine learning, both as a protocol for experiments with agents (Vaccaro, 2026) and as a defence against tuning prompts, decoding parameters or output format until a desired result appears (Thomas et al., 2026). We record our use of it as method rather than contribution.

Limitations

Substrate. Every measurement here is on function-level Python tasks from two code benchmarks. Nothing in this paper establishes that the same ceiling, or the same explanation for it, holds for another language, for repository-scale changes, or for code whose correctness is not decided by an executable suite.

The ground truth is a test suite, and a test suite is not a specification. Hidden suites encode a reference implementation’s behaviour, including behaviour its prose never fixed. In the intervention study reported in §3.5, a generator shown an instance’s failing inputs inferred a behavioural rule that *contradicted* the hidden oracle on 23 of 32 instances. The study cannot say whether that reflects the specifications or the generator’s inference from inputs alone; either way it bounds what any recall figure measured this way can mean.

Interval coverage, measured only in part. The primary intervals are percentile bootstraps over problem clusters. For a *single rate* only, simulated under the design’s own cluster profiles and an assumed within-problem correlation, with 600 resamples rather than the 10,000 the analyses use (400 replicates; Monte Carlo standard error 0.01–0.02), a nominal 95% interval covers between 0.93 and 0.97 at the tested true rates from 0.10 to 0.60, and 0.80 at a rate of 0.03 on the 56-cluster stratum (§2.3), so small-rate intervals are narrower than they claim. An earlier version of this paragraph gave a range the source report had withdrawn. Two earlier constructions in this programme were found by review to cover 0.416 and 0.075 where 0.95 was claimed.

The validation instrument shares the failure mode under study. Reports were checked by a language model, which is what this paper measures a ceiling on. Studies cleared review at between three and twenty-one rounds, with the round before the last always returning substantive findings; measurements moved on review in three studies, and elsewhere the refusals concerned interpretation, where the instrument is weakest; and withholding our conclusions from one pass produced six substantive findings that six briefed passes had missed. The defect class it leaves open is the one that looks right to a language model, and we do not

know what remains in it.

The reviewer is one of the measured families. `gpt-6-astra` reviewed every report in this programme and is also an auditor route whose ranking against the others the paper reports. For claims about `astra`’s relative performance the validation is a self-assessment, and those comparisons should be treated as bracketed until a model that audited nothing here has checked them.

No rater from outside the project. Every human judgement in this work is the first author’s. A third model rating does not discharge that, and the design that asks for an outside human has not been run.

Budget. Reading depth K , the number of arms, and the size of every intervention population were set by cost rather than by a power calculation, and no sample size was justified in advance by a minimum-detectable-effect argument.

Reproducibility and claims

The repository at <https://github.com/dongzhaoh321418-lab/audit-ceiling-paper> carries the claims ledger (CLAIMS.md), with the claims the evidence does not support; the results report of every study the paper draws on; every registered study’s preregistration and numbered amendments; every independent review of those reports, verbatim and including refusals; and the derived records — counts, labels, hashes and intervals — that the figures and Table 1 are drawn from. A single command, `reproduce.sh`, regenerates Figures 1–3 and Table 1 from those records and checks that each is byte-identical to the copy the manuscript compiles, with no model calls, no network access and no path outside the checkout; it was run from a fresh clone with the home directory unavailable to establish that. The repository’s README maps each claim to its report, its reviews and its records. The corpora are HumanEval (MIT) and HumanEval+ and MBPP+ (Apache-2.0), which permit redistribution. What a replication cannot do from this repository alone is regenerate the records themselves. With the archive of raw readings, which is kept separately, they regenerate deterministically through the analysis code in the auditor harness, a separate repository named in the provenance file; without it, only by re-running the studies against model endpoints whose outputs are not deterministic. The repository is private until the paper is released.

References

Brown, B., Juravsky, J., Ehrlich, R., Clark, R., Le, Q. V., Ré, C., and Mirhoseini, A. Large language monkeys: Scaling inference compute with repeated sampling. *arXiv*

preprint arXiv:2407.21787, 2024.

Chen, W.-L., Wei, Z., Zhu, X., Feng, S., and Meng, Y. Do LLM evaluators prefer themselves for a reason? *arXiv preprint arXiv:2504.03846*, 2025.

Huang, J., Chen, X., Mishra, S., Zheng, H. S., Yu, A. W., Song, X., and Zhou, D. Large language models cannot self-correct reasoning yet. *arXiv preprint arXiv:2310.01798*, 2024.

Konstantinou, M., Degiovanni, R., and Papadakis, M. Do LLMs generate test oracles that capture the actual or the expected program behaviour? *arXiv preprint arXiv:2410.21136*, 2024.

Liu, J., Xia, C. S., Wang, Y., and Zhang, L. Is your code generated by ChatGPT really correct? rigorous evaluation of large language models for code generation. In *Advances in Neural Information Processing Systems*, 2023. *arXiv:2305.01210*.

McAleese, N., Pokorny, R. M., Uribe, J. F. C., Nitishinskaya, E., Trebacz, M., and Leike, J. LLM critics help catch LLM bugs. *arXiv preprint arXiv:2407.00215*, 2024.

Richter, C. and Papadakis, M. Underspecification does not imply incoherence: The risks of semantic collapse in coding models. *arXiv preprint arXiv:2607.01953*, 2026.

Taherkhani, H., DaghighFarsoodeh, A., Chowdhury, M., Pham, H. V., and Hemmati, H. Consistency meets verification: Enhancing test generation quality in large language models without ground-truth solutions. *arXiv preprint arXiv:2602.10522*, 2026.

Tan, H., Sun, F., Liu, S., Su, D., Cao, Q., Chen, X., Wang, J., Cai, X., Wang, Y., Shen, H., and Cheng, X. Too consistent to detect: A study of self-consistent errors in LLMs. *arXiv preprint arXiv:2505.17656*, 2025.

Thomas, M., Gligorić, K., and Shah, N. B. Mitigating LLM-based p-hacking by preregistering for the next LLM. *arXiv preprint arXiv:2606.27687*, 2026.

Vaccaro, M. Preregistration for experiments with AI agents. *arXiv preprint arXiv:2606.11217*, 2026.

Xiong, W., Guo, Y., and Chen, H. The program testing ability of large language models for code. *arXiv preprint arXiv:2310.05727*, 2023.

Xu, Q., Wang, G., Briand, L., and Liu, K. Hallucination to consensus: Multi-agent LLMs for end-to-end JUnit test generation. *arXiv preprint arXiv:2506.02943*, 2025.

Zeng, Z., Shi, R., Han, K., Li, Y., Sun, K., Wang, Y., Yu, Z., Xie, R., Ye, W., and Zhang, S. SWR-Bench: Assessing LLM performance in real-world code review comment generation. *arXiv preprint arXiv:2509.01494*, 2025.